

Capstone Project – Chicago Crime Data

I. Introduction

The Chicago Data Portal publishes a dataset reflecting reported incidents of crime that occurred in the City of Chicago from 2001 to August 2018.

In this project, I have examined the data, conducted data wrangling, utilized different charts identifying general trends for the crime cases in Chicago, and applied inferential statistics techniques to the data. Further, I conducted broader correlation analyses between the crime cases and demographic factors such as population, education, unemployment etc. In the last part of the project, I applied some supervised machine learning techniques to the data and built a predictive model for the crime rate per community area and the relevant demographic factor. In addition, I also utilized an unsupervised machine learning technique to label community areas in Chicago and identify a cluster having the highest average crime rate.

The Chicago city government would be the clients for this project. They can use my analysis report to understand the impact of demographic factors on crime rates, and which factors play critical roles in crime rates. They can also use the built model to identify the crime rate for a community area, and focus on a group of community areas with the highest average crime rate. Subsequently, they can mobilize resources more efficiently and effectively and improve relevant areas in reducing the crime rate for a community area.

II. Data Sources

1. The data is available for download from the Chicago Data Portal as follows:

<https://data.cityofchicago.org/Public-Safety/Crimes-2001-to-present/ijzp-q8t2>

2. Chicago community area data covering population, education, income etc.

<https://datahub.cmap.illinois.gov/dataset/1d2dd970-f0a6-4736-96a1-3cae431f5e4/resource/8c4e096e-c90c-4bef-9cf1-9028d094296e/download/ReferenceCCA20112015.csv>

III. Completed Work

1. Data Wrangling:

Overall, the data is quite clean. With further exploration, some cleanups are conducted. The Jupyter Notebook is available in Github:

<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20Data%20Wrangling.ipynb>

- i. Remove rows containing at least one null value

The total number of cases is over 6.7 million, and the number of cases that has at least one NULL value is about 0.7 mil. Majority of the NULL values happen for the fields 'Ward' and 'Community Area'. Since the latter is only about 10% of the total cases, they are removed from the whole dataset.

- ii. Format date

- a. Some column names such as 'Community Area' include a space. I replace space with '_'.
- b. The fields of 'District', 'Ward', and 'Community_Area' are float64 data. Actually, they are just numeric codes contain no decimal point. I convert them to the int64 format.
- c. I changed the fields of 'IUCR', 'Primary_Type', and 'FBI_Code' to the categorical variables.
- d. I converted the 'Date' from the string format to the date-time format. Since the column 'Year' comes from the 'Date', at this moment, I remove it.

- iii. Remove duplicates

The field 'ID' is unique, but the field 'Case_Number' included 378 duplicates. I remove them from the dataset.

- iv. Remove outliers

I use the scatter plot to examine 'X_Coordinate' and 'Y_Coordinate', and find 107 outliers. I remove them from the dataset.

- v. Save the clean file

The dataset cleaned up gets saved to a file for data analysis.

2. Data Story Telling:

I utilized different plots to identify patterns of the crime cases. The Jupyter Notebook is available in Github:

<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20Data%20Storytelling.ipynb>

Below is the list of major trends and patterns identified

- i. The total number of crime cases decreased over the years from 2003 to 2015. However, the total number of crime cases stays flat between 2015 and 2018.
- ii. The number of crime cases drops to the lowest point around January or February each year, and the number of crime cases rises to the highest point around July each year.

- iii. The number of cases is the lowest between 4:00 – 6:00am. The peak hours for the largest number of crime cases varies with primary types of crime cases.
- iv. The top three primary types of crime cases are THEFT, BATTERY and CRIMINAL DAMAGE.
- v. The Community Area #25 has the largest number of crime cases.

3. Inferential Statistics Techniques:

The inferential statistics for the crime cases were also conducted by joining the 2011-2015 education data *** for Chicago.

- i. Below is a table for Pearson correlation coefficient between the crime rate and a feature for Community Areas

| Feature Name in Jupyter Book | Actual Feature Name | Pearson Correlation Coefficient between the Feature and Crime Rate | Notes |
|------------------------------|------------------------------------|--|--------------------------------|
| MED_AGE | Median age | -0.25821 | |
| HS_RATE | High school diploma or higher rate | -0.29572 | Based on Population 19 or over |
| BACH_RATE | Bachelor degree or higher rate | -0.46238 | Based on Population 19 or over |
| UNEMP_RATE | Unemployment rate | 0.80690 | |
| MEDINC | Median income | -0.67329 | |
| OWN_OCC_HU_RATE | House/Apartment owning rate | -0.61656 | |

- ii. The Unemployment rate has very strong correlation with the Crime Rate. The rest of factors have more or less negative correlations with the Crime Rate.
- iii. Some features have some correlations to some degree. For example, the Bachelor degree or higher rate and the High school diploma or higher rate have some overlaps.

4. Supervised Machine Learning - Linear Regressions

- i. Linear Regression

a) Linear Regression

Use the six features listed in the table in 3.i, conducted Linear Regression. The R-squared is 0.75986, which means that about 75.99% of the Crime Rate change can be explained by these six features. F-statistic also indicates that the model built is statistically significant overall. However, through the 5-fold cross validation, the cross validation score varies quite largely. The average score is 0.5315.

b) Ridge Regression

With the technique, split the dataset into two parts: training and test. The predication score for the test set is 0.78409.

With the 5-fold cross validation, the cross validation score varies in a much narrower range than that for the previous Linear Regression. The average score is 0.58675. This regression technique yields better results (model).

c) Lasso Regression

With this technique, we can conclude that the Unemployment Rate is the dominant feature in determining the Crime Rate, and the House/Apartment Rate is the second important feature impacting the Crime Rate.

5. Unsupervised Machine Learning – KMeans Clustering

After producing the graph between Inertia and the number of Clusters for KMeans, and conducting several tries, I picked the 7 as the optimize clusters. With this technique, I identified a group of Community Areas having the highest average Crime Rate. I also generated a graph for the Crime Rate and features on this group against the average values for all of Community Areas.

V. Deliverables

By completion of this project, I will have the following deliverables:

- Jupyter books for data analysis and machine learning
 - Data Wrangling:
<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20Data%20Wrangling.ipynb>
 - Data Story Telling:
<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20Data%20Storytelling.ipynb>
 - Inferential Statistics Techniques:
<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20MachineLearning.ipynb>
 - Supervised Machine Learning - Linear Regressions:
<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20MachineLearning.ipynb>
 - Unsupervised Machine Learning – KMeans Clustering:
<https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeData%20-%20MachineLearning.ipynb>
- A document describe the approach to be adopted (this document)
- A slide deck for the whole project and conclusion:
https://github.com/DavidFlanders/Capstone/blob/master/ChicagoCrimeDataset_Project_UserPresentation.pdf
-

VI. Appendix, Description of Chicago Crime Dataset

| Column Name | Description | Type |
|----------------------|---|-------------|
| ID | Unique identifier for the record. | Number |
| Case Number | The Chicago Police Department RD Number (Records Division Number), which is unique to the incident. | Plain Text |
| Date | Date when the incident occurred. This is sometimes a best estimate. | Date & Time |
| Block | The partially redacted address where the incident occurred, placing it on the same block as the actual address. | Plain Text |
| IUCR | The Illinois Uniform Crime Reporting code. This is directly linked to the Primary Type and Description. See the list of IUCR codes at https://data.cityofchicago.org/d/c7ck-438e . | Plain Text |
| Primary Type | The primary description of the IUCR code. | Plain Text |
| Description | The secondary description of the IUCR code, a subcategory of the primary description. | Plain Text |
| Location Description | Description of the location where the incident occurred. | Plain Text |
| Arrest | Indicates whether an arrest was made. | Checkbox |
| Domestic | Indicates whether the incident was domestic-related as defined by the Illinois Domestic Violence Act. | Checkbox |

| | | |
|----------------|--|-------------|
| Beat | Indicates the beat where the incident occurred. A beat is the smallest police geographic area – each beat has a dedicated police beat car. Three to five beats make up a police sector, and three sectors make up a police district. The Chicago Police Department has 22 police districts. See the beats at https://data.cityofchicago.org/d/aerh-rz74 . | Plain Text |
| District | Indicates the police district where the incident occurred. See the districts at https://data.cityofchicago.org/d/fthy-xz3r . | Plain Text |
| Ward | The ward (City Council district) where the incident occurred. See the wards at https://data.cityofchicago.org/d/sp34-6z76 . | Number |
| Community Area | Indicates the community area where the incident occurred. Chicago has 77 community areas. See the community areas at https://data.cityofchicago.org/d/cauq-8yn6 . | Plain Text |
| FBI Code | Indicates the crime classification as outlined in the FBI's National Incident-Based Reporting System (NIBRS). See the Chicago Police Department listing of these classifications at http://gis.chicagopolice.org/clearmap_crime_sums/crime_types.html . | Plain Text |
| X Coordinate | The x coordinate of the location where the incident occurred in State Plane Illinois East NAD 1983 projection. This location is shifted from the actual location for partial redaction but falls on the same block. | Number |
| Y Coordinate | The y coordinate of the location where the incident occurred in State Plane Illinois East NAD 1983 projection. This location is shifted from the actual location for partial redaction but falls on the same block. | Number |
| Year | Year the incident occurred. | Number |
| Updated On | Date and time the record was last updated. | Date & Time |

| | | |
|-----------|---|--------|
| Latitude | The latitude of the location where the incident occurred. This location is shifted from the actual location for partial redaction but falls on the same block. | Number |
| Longitude | The longitude of the location where the incident occurred. This location is shifted from the actual location for partial redaction but falls on the same block. | Number |
| Location | The location where the incident occurred in a format that allows for creation of maps and other geographic operations on this data portal. This location is shifted from the actual location for partial redaction but falls on the same block. | |